

Adventures in Extension Packaging

**David Wheeler
PGXN, Tembo**

HELLO!

I'M DAVID WHEELER

Principle Architect, Tembo

PostgreSQL user since 2000

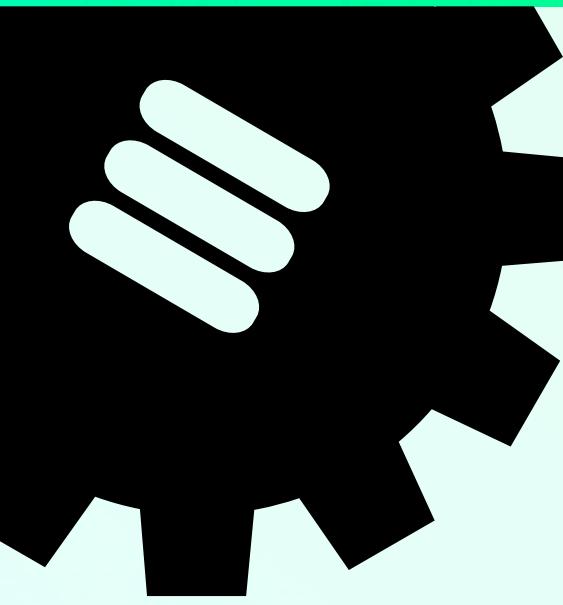
PostgreSQL contributor since 2008

pgTAP: Unit testing for PostgreSQL

Sqitch: Database change management

PGXN: PostgreSQL Extension network





PGXN

pgxn.org

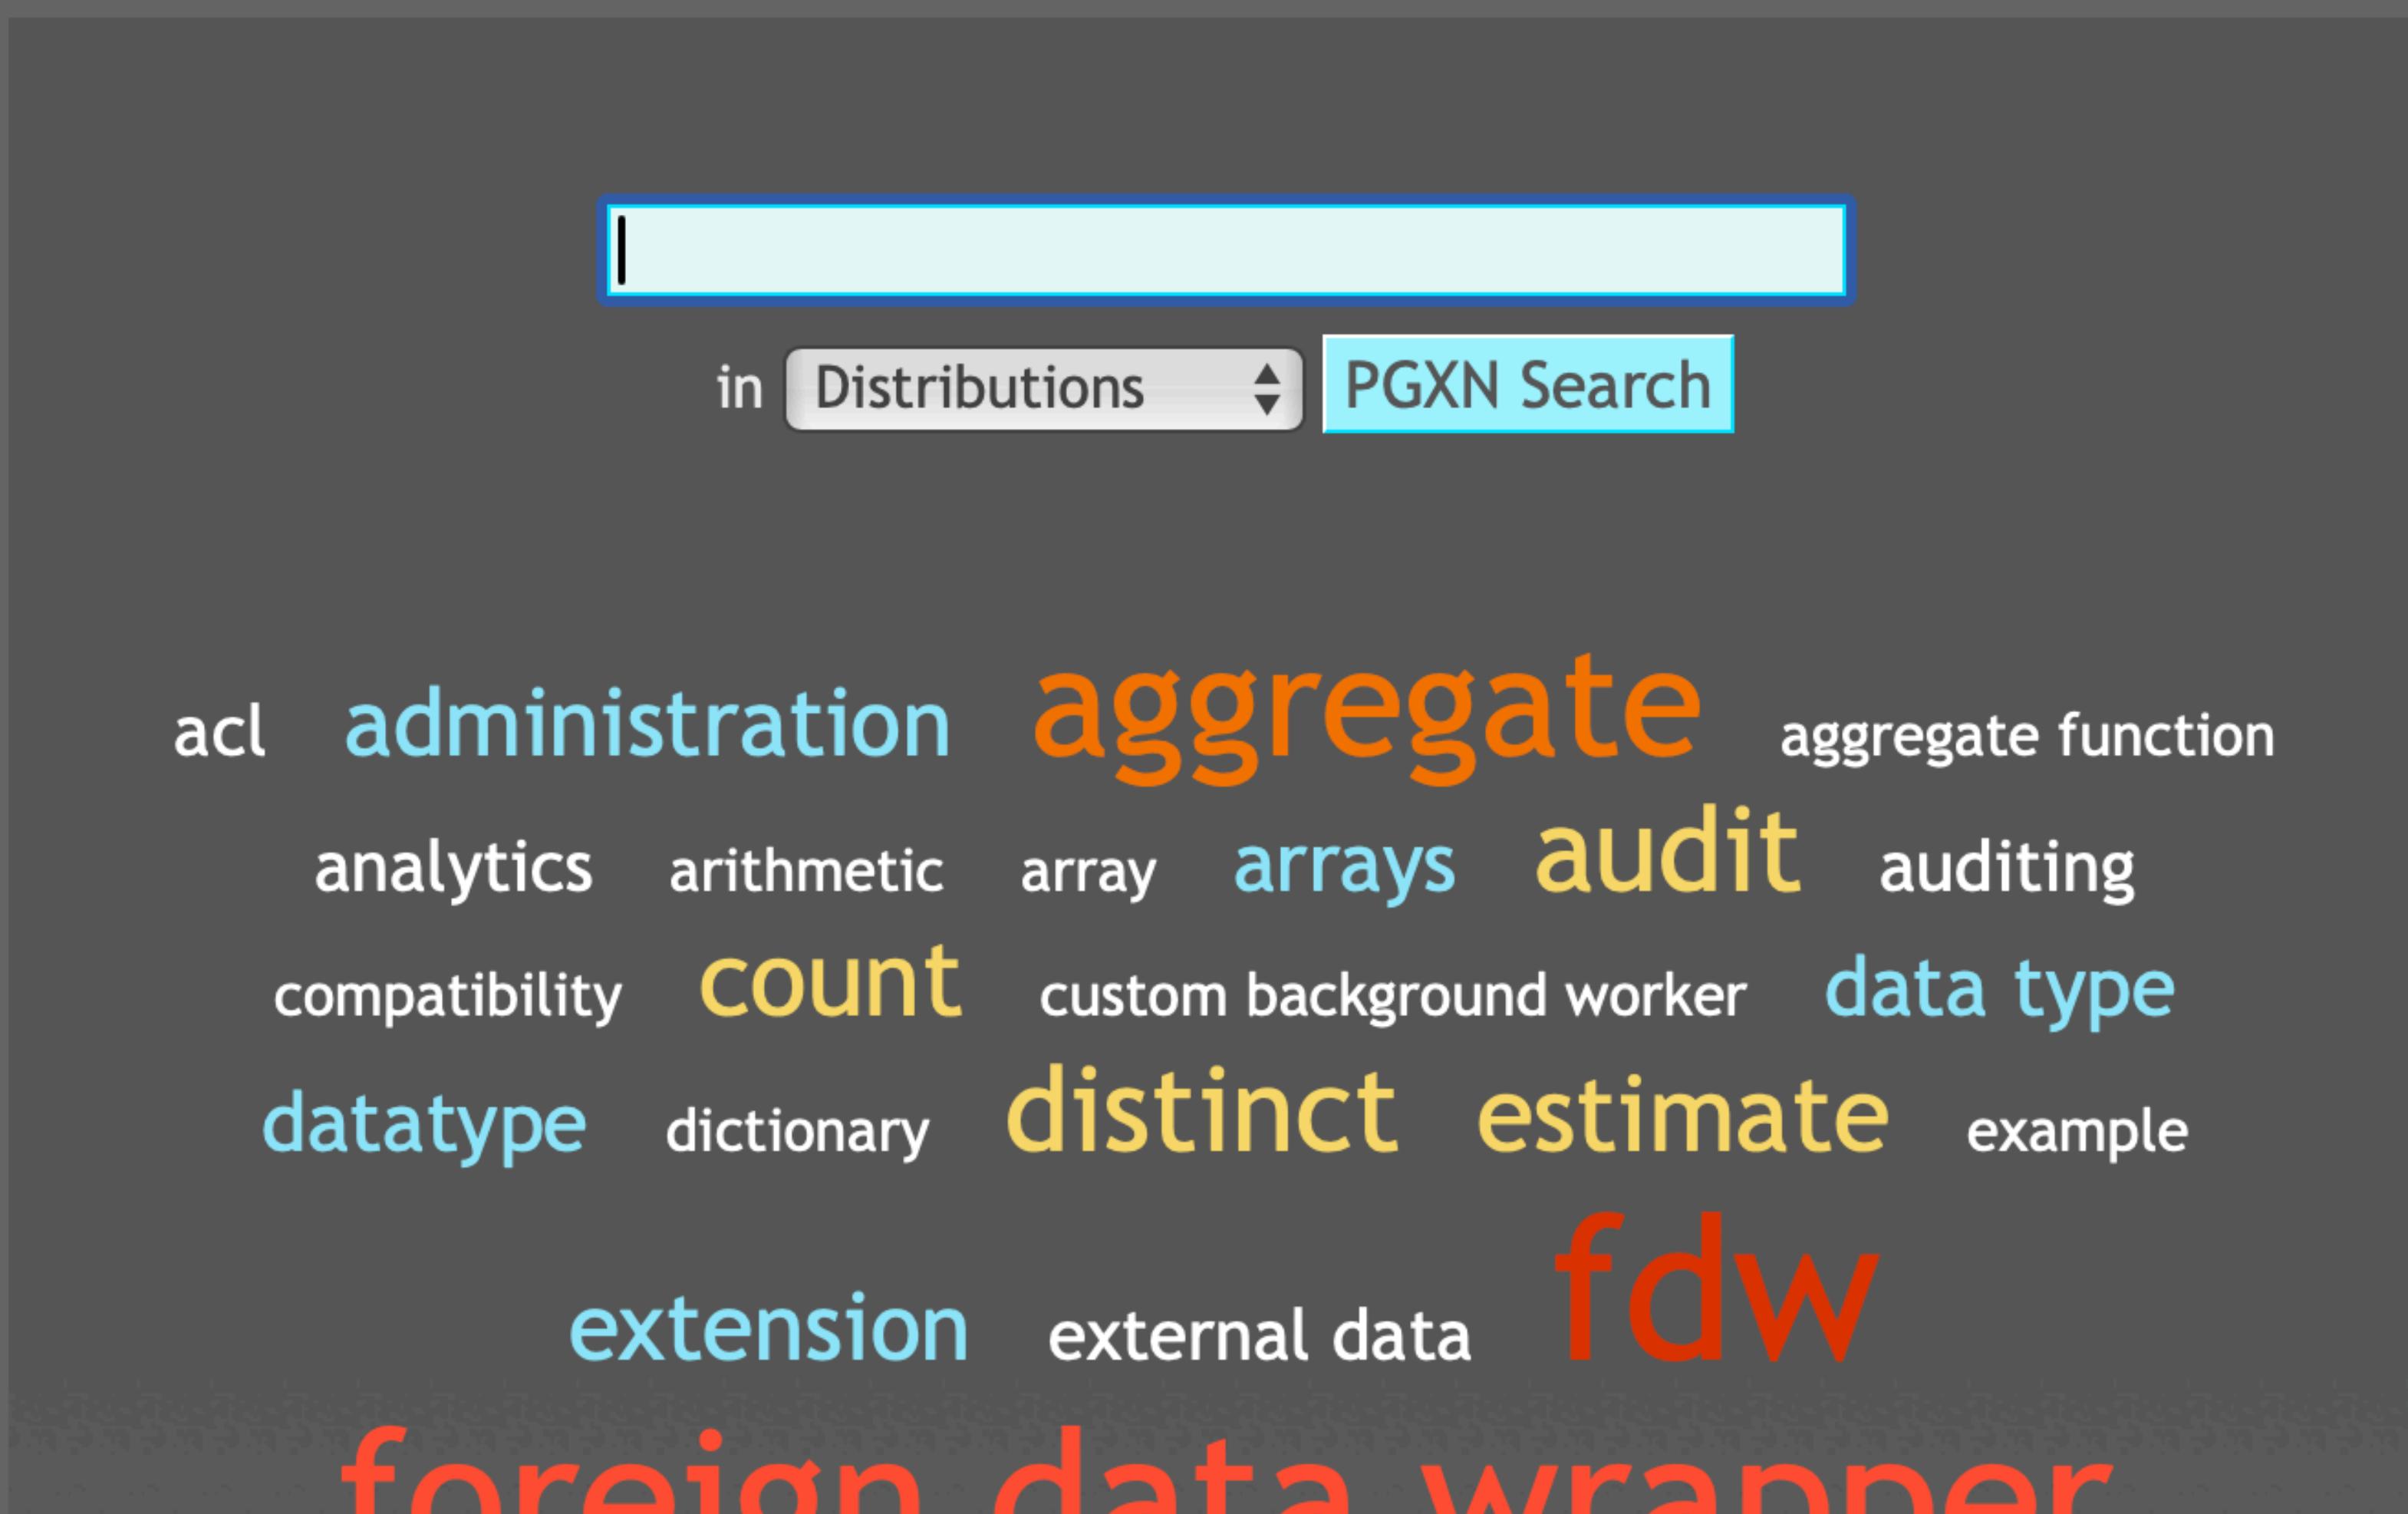
Established 2011

Source code distribution network

PostgreSQL extensions and tools

Search, documentation, tags

CLI by Daniele Verazzo

[users](#) [tags](#) [recent](#)

The screenshot shows a search interface for PostgreSQL extensions. At the top, there is a search bar with a placeholder 'Search' and a dropdown menu set to 'in Distributions'. Below the search bar is a grid of extension names. The extensions listed are: acl, administration, aggregate, analytics, arithmetic, array, arrays, audit, compatibility, COUNT, custom background worker, datatype, dictionary, distinct, estimate, extension, external data, fdw, foreign data wrapper, uint128, and uint128.

acl	administration	aggregate	aggregate function		
analytics	arithmetic	array	arrays	audit	auditing
compatibility	COUNT	custom background worker	datatype		
datatype	dictionary	distinct	estimate	example	
		fdw			
		foreign data wrapper			
		uint128	1.0.1		

PGXN, the PostgreSQL Extension network, is a central distribution system for open-source PostgreSQL extension libraries.

Recent Releases

[pg_task 2.1.16](#)
PostgreSQL and Greenplum job scheduler pg_task allows to execute any sql command at any specific time at background asynchronously

[pg_mustach 1.0.16](#)
PostgreSQL mustach

EXTENSIONS

422 Extensions

398 Distributions

2490 Releases

468 Users

2 Mirrors

But...

jоelоnsql/PostgreSQL-EXTENSIONS.md: 1,190

**It's not
enough.**



THE PGXN v2 VISION

PGXN V2 FEATURES

Canonical source registry

Search & discover all extensions

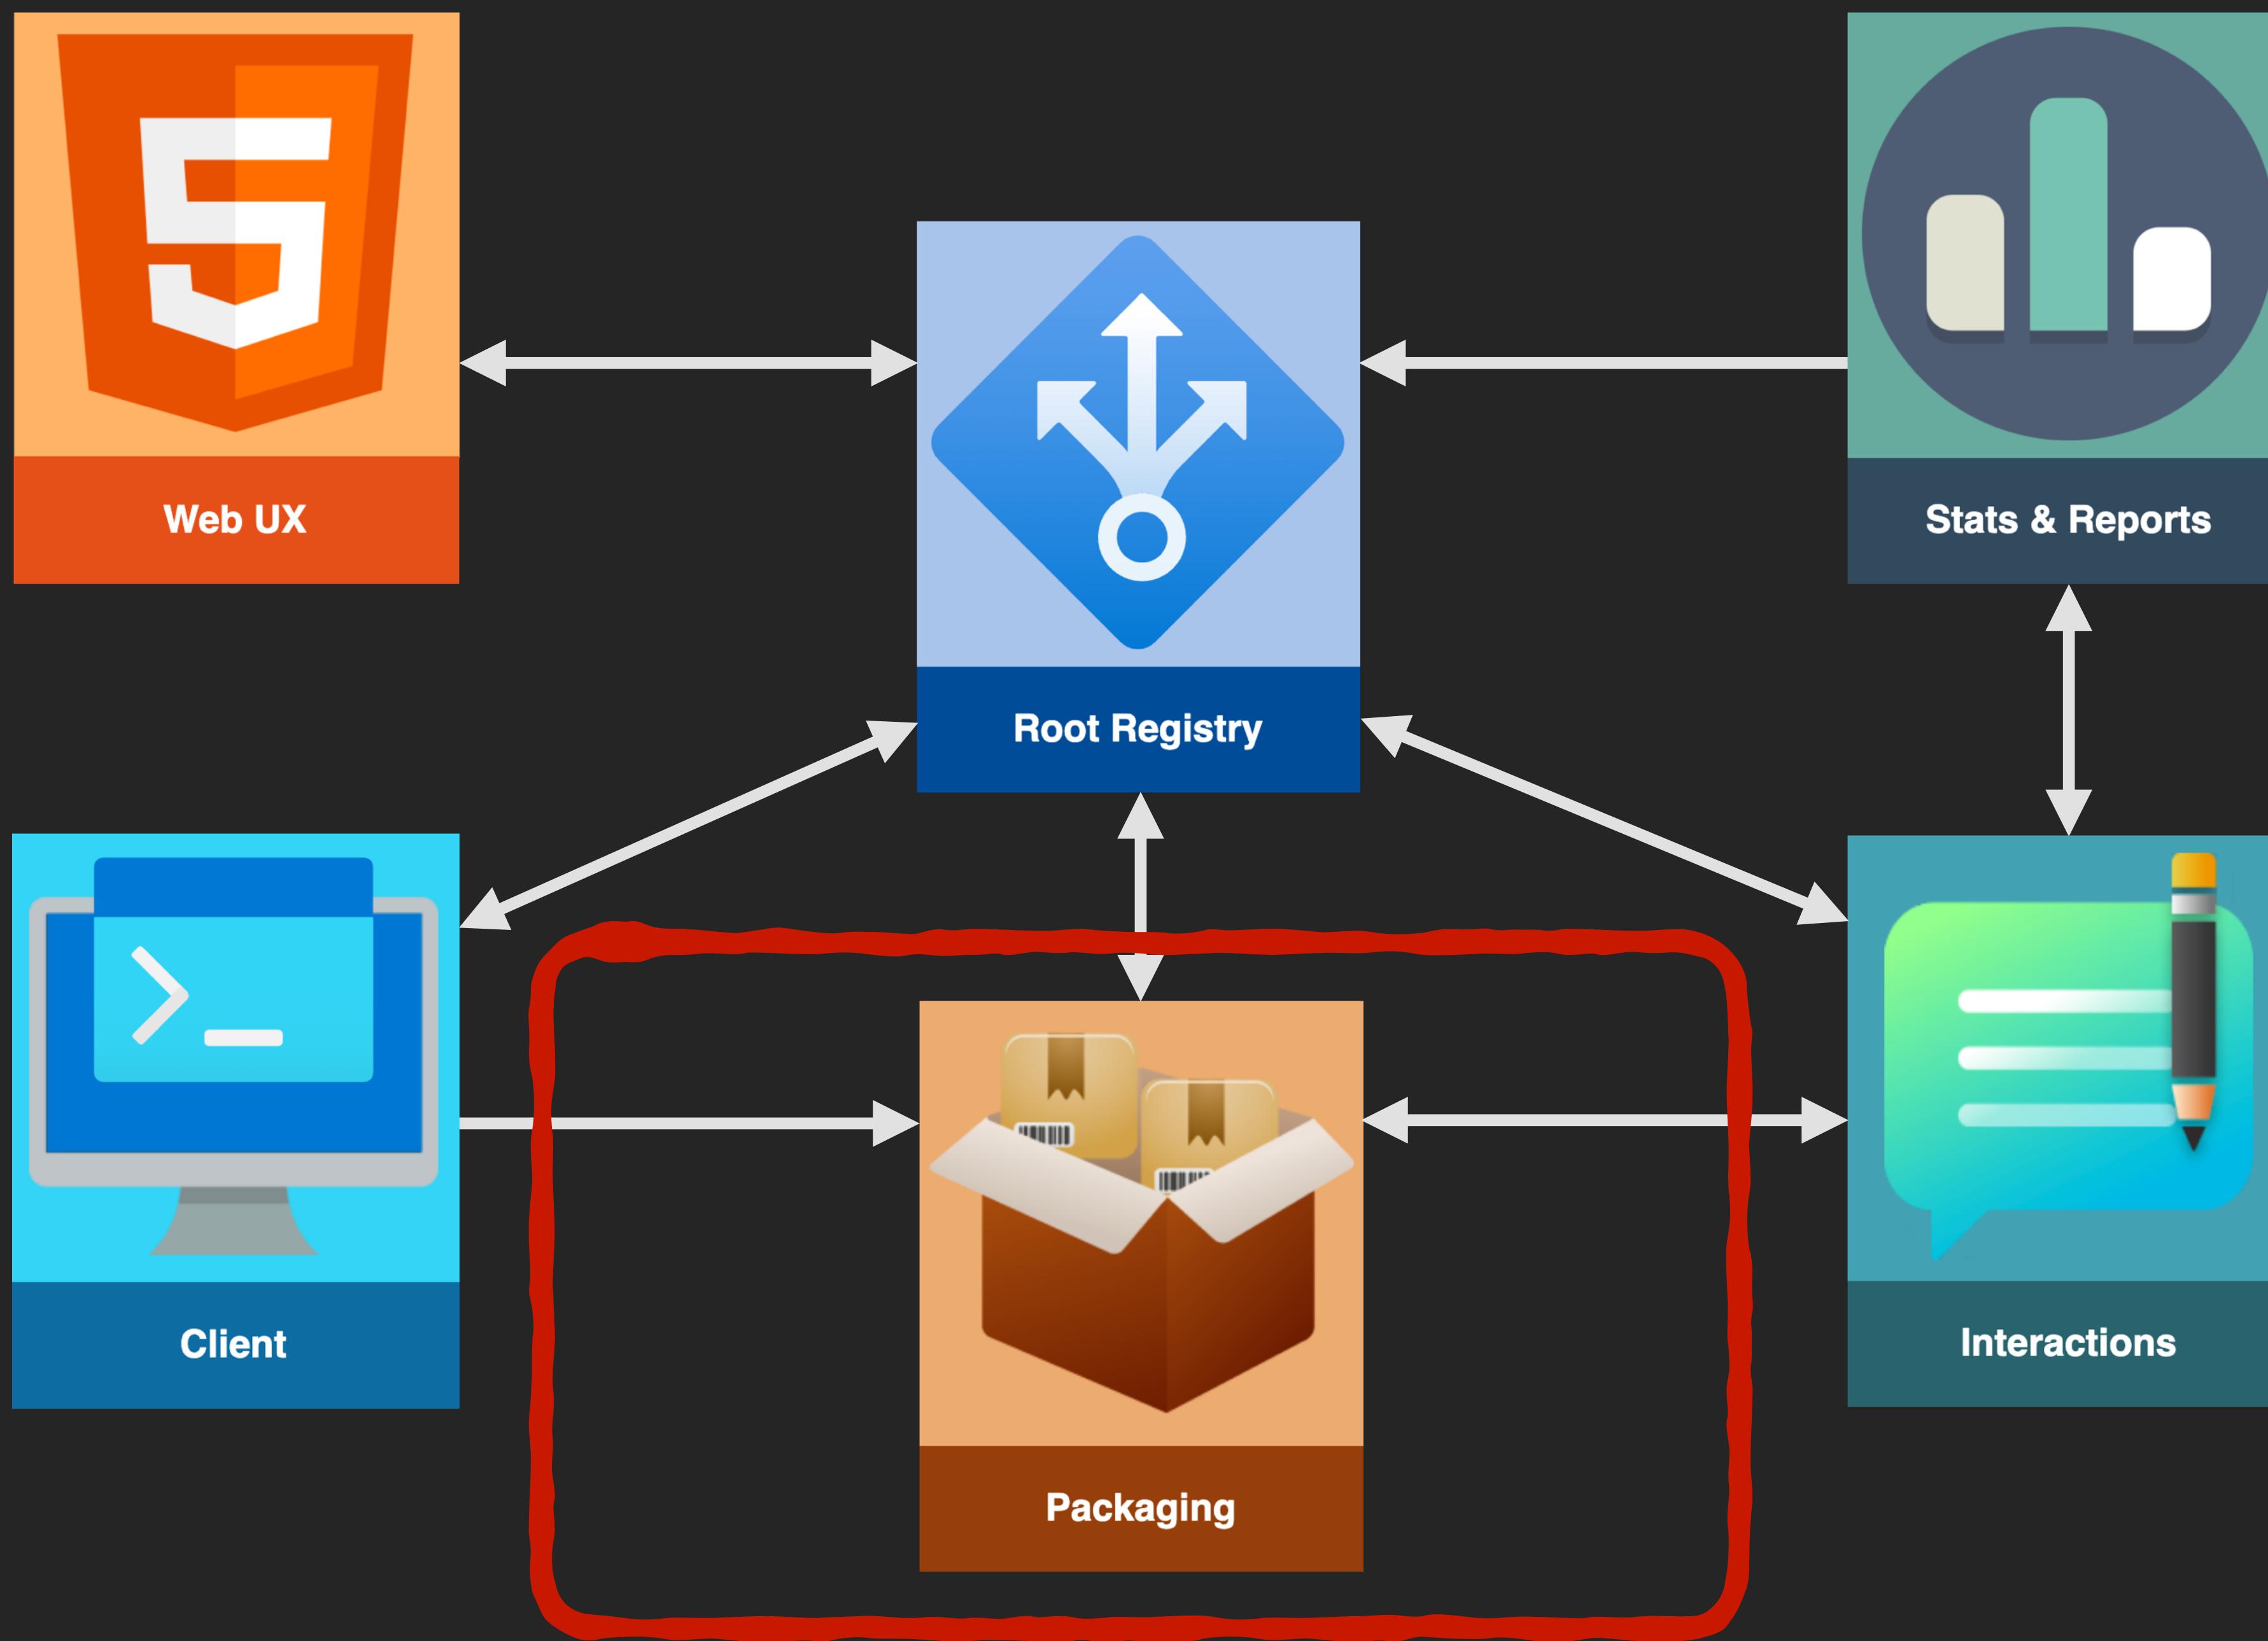
Streaming & federation

Quality metrics & reports

Binary packaging & distribution

Multi-platform & multi-version

Simple installation & package management



METADATA & PACKAGING RFCS

github.com/pgxn/rfcs/pulls

Started developing RFCs

Meta Spec v2

JWS Release Signing

Binary Distribution Format

Feedback appreciated!

Goals:

META.json-driven packaging

OCI distribution

META SPEC v2

github.com/pgxn/rfcs/pulls

Indexing: Contents

Extensions, modules and apps

Discovery: Classifications

Tags and categories

Dependencies:

Postgres, build tooling, other extensions, external

**Also:
build,
develop**

```
{  
  "dependencies": [  
    "postgres": {  
      "version": ">= 12.0, < 17.0",  
      "with": [ "xml", "uuid", "perl" ]  
    },  
    "pgrx": "pgrx",  
    "packages": {  
      "configure": {  
        "pkg": "cargo/cargo-pgrx": "==0.14.2",  
        "pkg": "generic/cmake": 0  
      }  
    },  
    "test": {  
      "requires": {  
        "pkg": "postgres/pg_regress": 0,  
        "pkg": "pgx/pgx": "1.1.0"  
      }  
    },  
    "run": {  
      "requires": {  
        "pkg": "pgxn/hostname": 0  
      }  
    }  
  }  
}
```

purl:
Package URL

```
{  
  "variations": [  
    {  
      "where": {  
        "platforms": ["linux"]  
      },  
      "dependencies": {  
        "packages": {  
          "run": {  
            "recommends": {  
              "pkg:pypi/auto-gptq": 0,  
              "pkg:pypi/xformers": 0  
            }  
          }  
        }  
      }  
    }  
  ]  
}
```

IMAGINE THIS

Source code release on PGXN

Interactions service calls registry Webhooks

Yum, APT, Homebrew, Pigsty, DBDev, etc.

Registries automatically pull new release

Install/configure dependencies from spec

Build, test, package, release

Report results back to PGXN

CLI to install packages for each platform

Package Management

› pgxn search json

rating	package	version	date	
★★	hadi/json_fdw	1.0.0	2013-11-18	Foreign Data
★★★★★	theory/jsonschema	0.1.0	2024-04-30	JSON Schema
★★★	jma/yamltodb	0.10.0	2022-11-08	Database ch
★★	shanekilkelly/bedquilt	2.1.0	2016-09-12	A json docu
★★★	furstenheim/is_jsonb_valid	0.0.1	2017-08-17	JSONB valid
★★★	postgrespro/jsonb_schema	1.0.0	2020-08-24	Decouple js
★★	kungfury/jsoncdc	0.1.0	2017-11-23	Translates

›

Package Management

```
> pgxn install -v theory/jsonschema
```

```
Installing theory/jsonschema 0.1.0 into pgenv/16.3.  
Found theory-jsonschema-0.1.1+2.pg16+1.darwin-arm64.zip  
Unpacking theory-jsonschema-0.1.1+2.pg16+1.darwin-arm64.zip  
Installing lib/jsonschema.dylib in ~/.pgenv/pgsql-16.3/lib/  
Installing share/jsonschema-0.1.0sql in ~/.pgenv/pgsql-16.3/share/extension  
Installing share/jsonschema.control in ~/.pgenv/pgsql-16.3/share/extension  
Installing doc/jsonschema.md in ~/.pgenv/pgsql-16.3/share/doc/
```

```
Done!  
>
```

Package Management

```
> pgxn package ls
```

Packages installed in pgenv/16.3:

package	installed	latest	description
core/citext	1.6.0	✓	data type for case-insensitive ch
core/jsonb_plperl	1.0.0	✓	transform between jsonb and plperl
core/unaccent	1.1.0	✓	text search dictionary that remov
theory/jsonschema	0.1.0	✓	JSON Schema validation functions

```
>
```



UNIVERSAL PACKAGING & DISTRIBUTION

TRUNK PACKAGE FORMAT

File layout defined by the package format

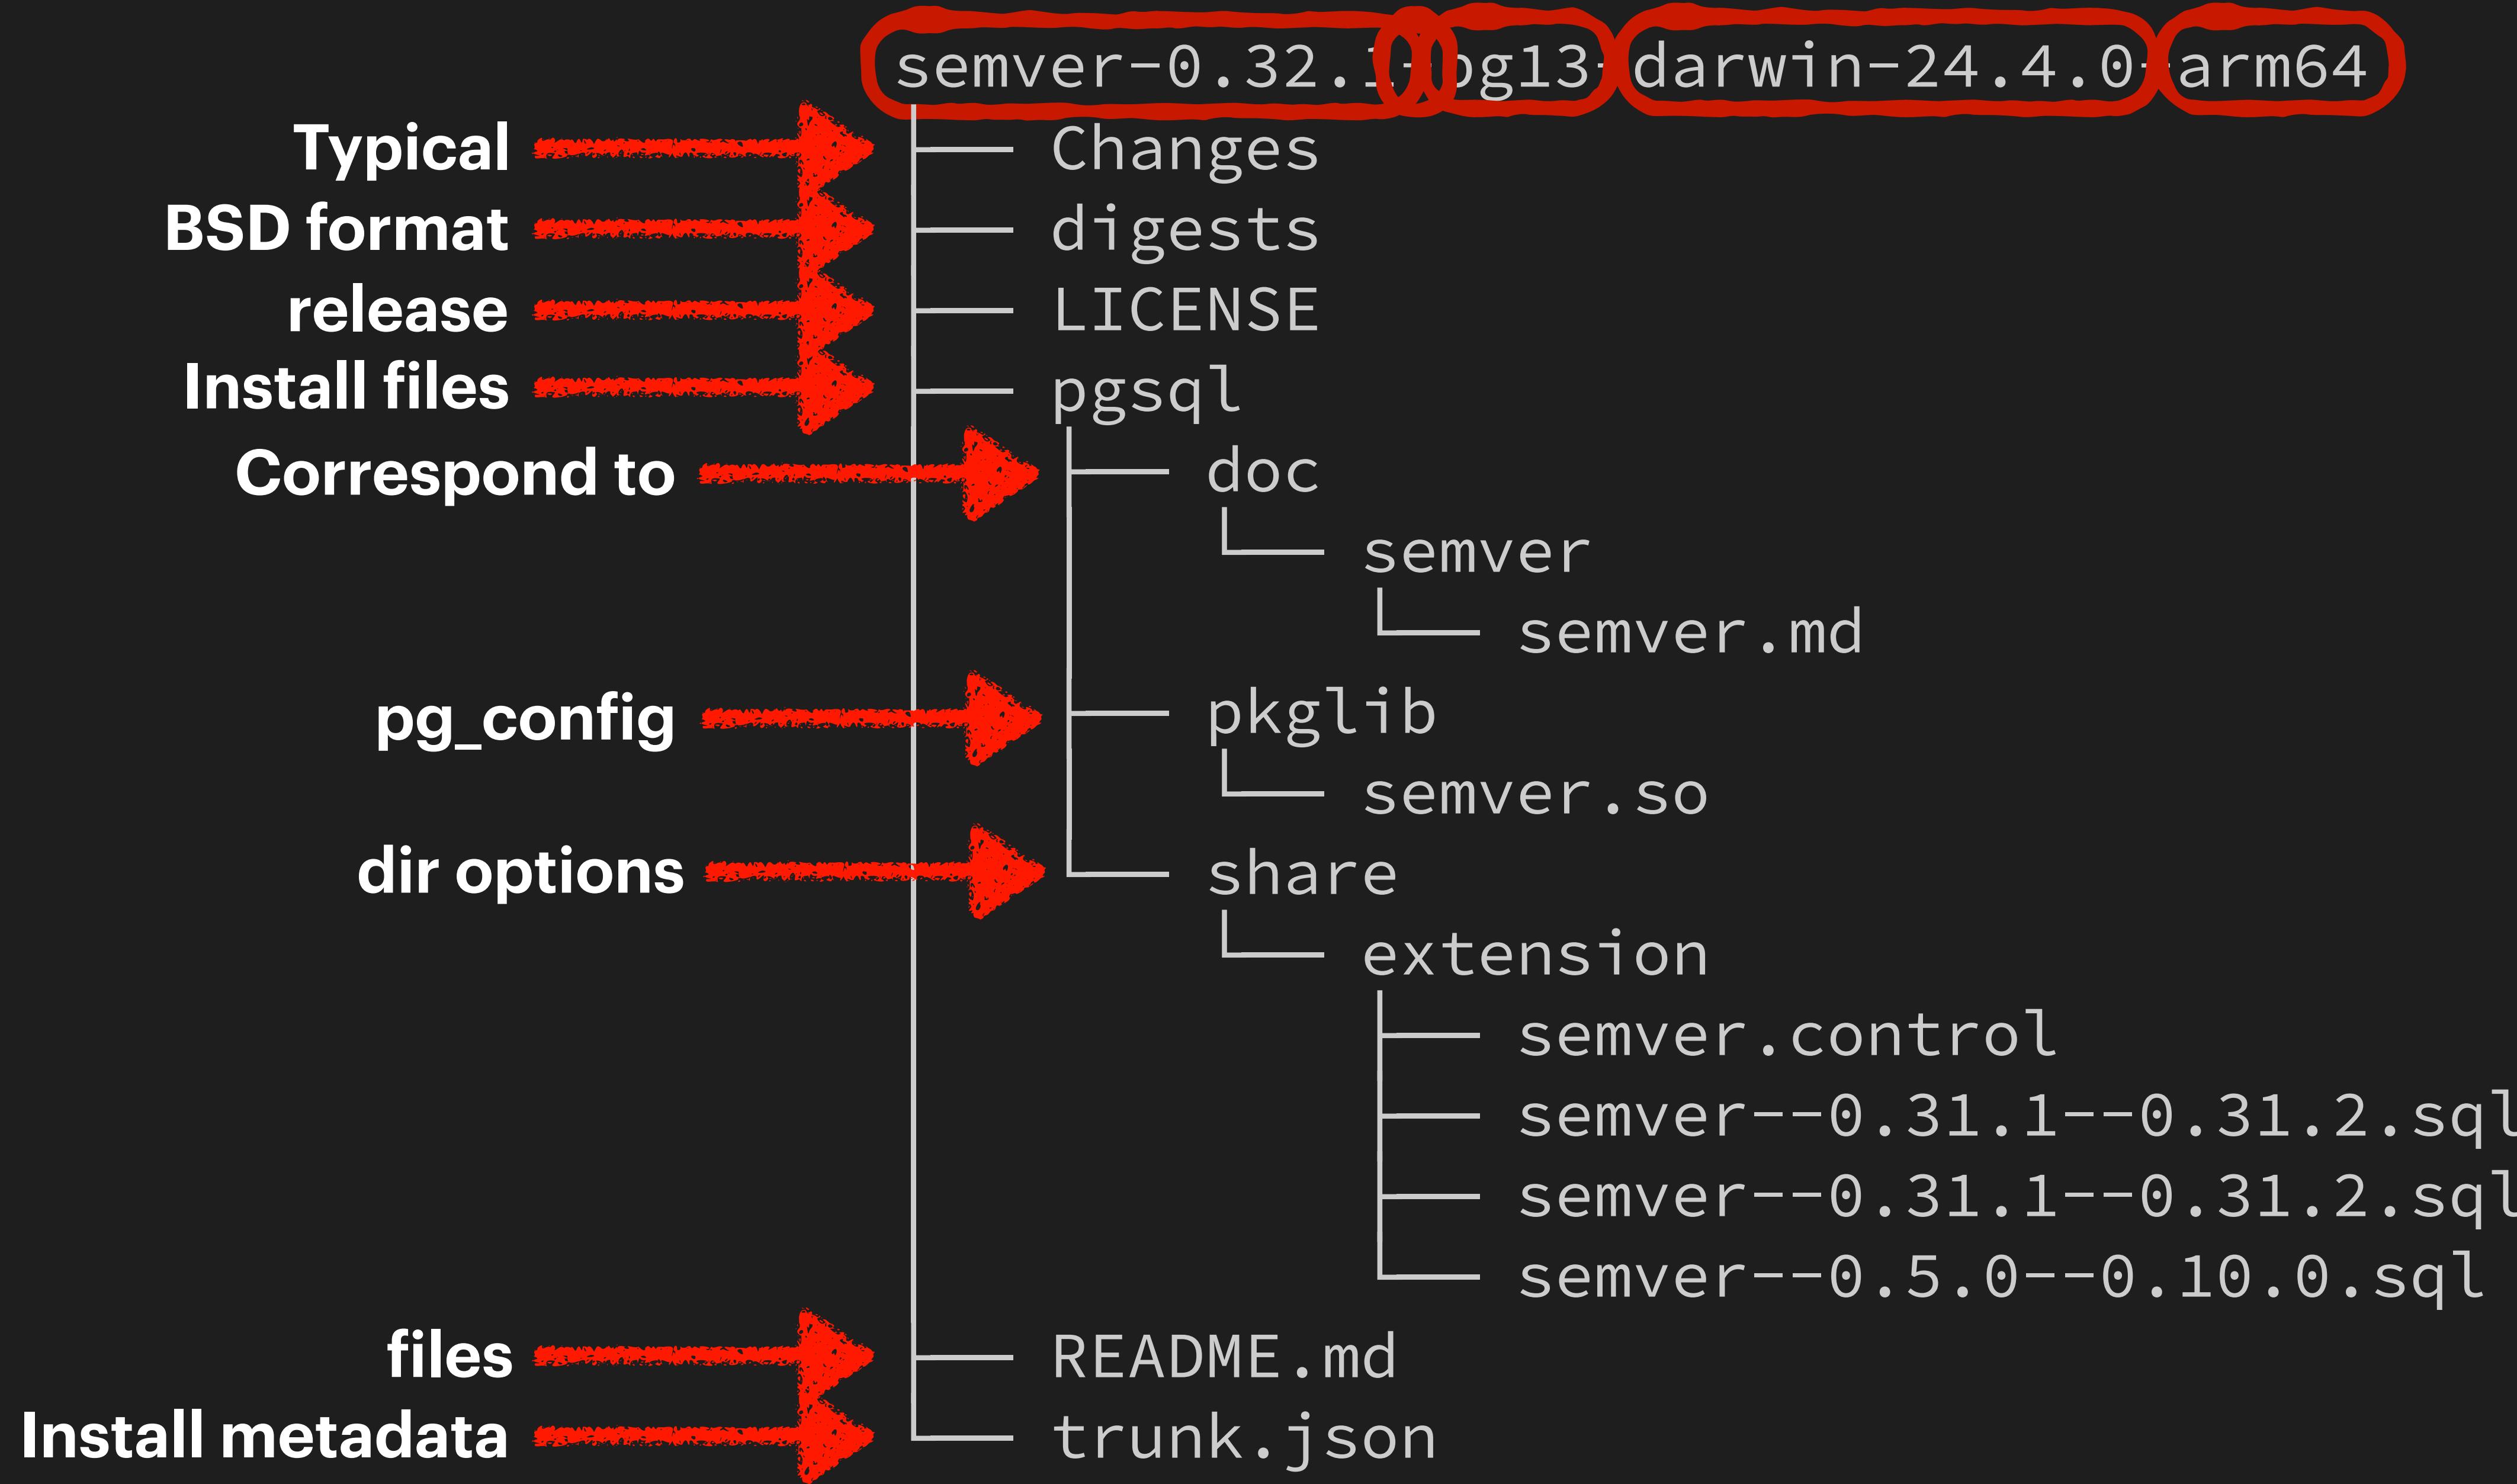
Not tied to OS layout

Differs from APT/RPM

Use absolute paths

Allows installation on any compatible Postgres

Regardless of Postgres install location



trunk.json

```
{  
  "trunk": "0.1.0",  
  "package": {  
    "name": "semver",  
    "version": "0.32.1",  
    "language": "c",  
    "license": "PostgreSQL"  
  },  
  "postgres": {  
    "version": "17.4",  
    "major": "17",  
    "number": 170004,  
    "libs": "-lpgcommon -lpgport -lxmll2 -lssl -lcrypto -lz -lreadline -lm ",  
    "cppflags": "-I. -I./ -I/Users/david/.pgenv/pgsql-17.4/include/server -I/Users/david/.pgenv/",  
    "cflags": "-Wall -Wmissing-prototypes -Wpointer-arith -Wdeclaration-after-statement -Werror=v",  
    "ldflags": "-L/Users/david/.pgenv/pgsql-17.4/lib -isysroot /Library/Developer/CommandLineTool  
  },  
  "platform": {  
    "os": "darwin",  
    "version": "24.4.0",  
    "arch": "arm64"  
  }  
}
```

OCI DISTRIBUTION

Compatible with OCI artifact format

Just bundle files into an image

Index platforms in image index manifest

Automatically install correct binary

Can extend metadata for Postgres version

Plus any other platform metadata

Standard OCI Metadata

```
{  
  "org.opencontainers.image.created": "2024-06-20T18:07:24Z",  
  "org.opencontainers.image.licenses": "PostgreSQL",  
  "org.opencontainers.image.title": "semver",  
  "org.opencontainers.image.description": "A Postgres data type Semantic Versions.",  
  "org.opencontainers.image.source": "https://github.com/theory/pg-semver",  
  "org.opencontainers.image.vendor": "PGXN",  
  "org.opencontainers.image.ref.name": "0.32.1",  
  "org.opencontainers.image.version": "0.32.1",  
  "org.opencontainers.image.url": "https://github.com/theory/pg-semver"  
}
```

Standard Platform Metadata

```
{  
  "os": "darwin",  
  "os.version": "23.5.0",  
  "architecture": "arm64",  
  "created": "2025-05-14T18:07:24Z"  
}
```

```
{  
  "os": "linux",  
  "architecture": "amd64",  
  "created": "2025-05-14T18:07:24Z"  
}
```

Custom PGXN Metadata

```
{  
  "org.pgxn.trunk.pg.version": "17.4",  
  "org.pgxn.trunk.pg.major": "17",  
  "org.pgxn.trunk.pg.version_num": "170004",  
  "org.pgxn.trunk.version": "0.1.0"  
}
```



DEMO



THAT'S THE DREAM

A Postgres

Extension Registry

Get Trunk

Cargo

cargo install pg-trunk

Copy

Trunk is an open-source package installer and registry for PostgreSQL extensions.

Featured

7

Analytics

22

Auditing / Logging

8

All Extensions

Search 238 extensions



Search

PACKAGING AUTOMATION CHALLENGES

Custom configuration (e.g., Rust package, RUSTFLAGS)

Binary compatibility (e.g., libc)

Packaging variation (e.g., Yum vs. APT)

Dependency package availability

Conflicting dependencies (e.g., pg_duckdb vs duckdb_fdw)

**Specific pgrx versions
Patches**

**Inconsistent build pipelines
Separate releases for different Postgres versions (pgaudit, pg_hint_plan)**

Extra build steps (bson, timescaledb)

NOT INSURMOUNTABLE

Add additional build configuration

Balance against exploits

No commands!

Merge additional metadata

Auto-report failures upstream

Build binaries for each libc

Build binaries for each platform

Start focused, grow from there

Risk of
combinatorial
explosion?

**How I imagine Christoph and Devrim
right about now...**

AND THEN HE SAID...

**WE'LL PACKAGE EVERY EXTENSION
FOR EVERY PLATFORM!**



Or maybe...



AND THEN HE SAID...

**UPSTREAM MAINTAINERS WILL
FIX BUILD FAILURES!**



A CNPG SIDE QUEST

EXTENSIONS IN CNPG

CloudNativePG containers immutable

Extensions must be baked in

Postgres 18 adds extension_control_path

Kubernetes 1.33 adds image volumes (beta)

CNPG 1.26 adds extensions to immutable containers

CNPG With Extension

```
apiVersion: postgresql.cnpg.io/v1
kind: Cluster
metadata:
  name: postgresql-with-extensions
spec:
  instances: 1
  imageName: ghcr.io/cloudnative-pg/postgresql-trunk:18-devel
  postgresql:
    extensions:
      - name: pgvector
        image:
          reference: ghcr.io/cloudnative-pg/pgvector-18-testing:latest
  storage:
    storageClass: standard
    size: 1Gi
```



A red box highlights the 'ImageVolume resource' section of the configuration. This section contains the 'extensions' and 'image' fields for the 'pgvector' extension.

CNPG EXTENSION DEPLOYMENT

1. Triggers rolling update, replicas first
2. Each extension mounted as read-only volume

`/extensions/pgvector`

3. CNPG Updates GUCs

`extension_control_path = '$system:/extensions/pgvector/share'`

`dynamic_library_path = '$libdir:/extensions/pgvector/lib'`

CNPG CHALLENGES

This works!

Some limitations

Volume resolves at pod startup

Updating GUCs requires restart

Cannot use single prefix for all image volumes

Each operates like full server prefix

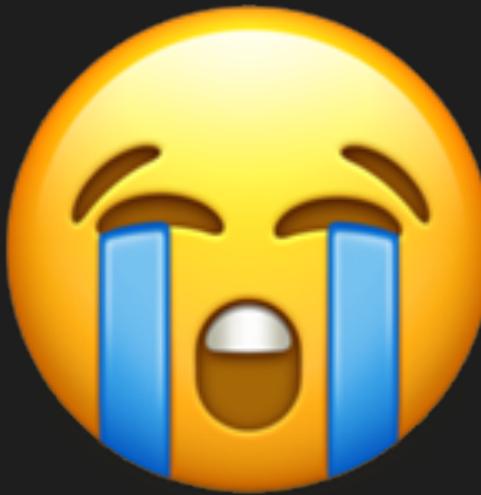
Multiple Extensions

```
extensions:
  - name: pgvector
    image:
      reference: ghcr.io/cloudnative-pg/pgvector-18-testing:latest
  - name: semver
    image:
      reference: ghcr.io/example/semver:0.40.0
  - name: auto_explain
    image:
      reference: ghcr.io/example/auto_explain:18
  - name: bloom
    image:
      reference: ghcr.io/example/bloom:18
  - name: postgis
    image:
      reference: ghcr.io/example/postgis:18
```

Multiple Extensions

```
extension_control_path = '$system:/extensions/pgvector/share:/  
extensions/semver/share:/extensions/auto_explain/share:/extensions/  
bloom/share:/extensions/postgis/share'
```

```
dynamic_library_path = '$libdir:/extensions/pgvector/lib:/  
extensions/semver/lib:/extensions/auto_explain/lib:/extensions/bloom/  
lib:/extensions/postgis/lib'
```



THE IDEAL

Add extensions by patching manifest

No rolling restarts

 **Requires ImageVolume without pod start**

No GUCs updated

Except shared_preload_libraries

 **Requires different file layout**

Would also simplify non-CNPG installs



RFC: EXTENSION PACKAGING & LOOKUP

IMAGINE THIS

Single search path GUC

No magic “postgresql” suffix

Each extension in eponymous directory

Maybe versioned?

Pre-defined file subdirectory names

Extension Search Path

```
extension_search_path =  
'$system:/extensions:/usr/local/extensions'
```

Extension Directory Layout

```
› ls -1 extensions
```

```
auto_explain
```

```
bloom
```

```
postgis
```

```
semver
```

```
vector
```

```
› tree extensions/semver
```

```
extensions/semver
```

```
├── doc
```

```
│   └── semver.md
```

```
└── lib
```

```
    └── semver.so
```

```
└── semver.control
```

```
└── sql
```

```
    ├── semver--0.31.0--0.31.1.sql
```

```
    ├── semver--0.31.1--0.31.2.sql
```

```
    ├── semver--0.31.2--0.32.0.sql
```

```
    └── semver--0.5.0--0.10.0.sql
```

Postgres looks for
directory name

Also: **bin, man,**
html, include,
locale

TRUNK FORMAT

Make trunk format identical to extension directory layout

Installation becomes simple:

Pull package from OCI registry

Validate signature & contents

Move entire directory to extension search path

CNPG AND TRUNK FORMAT

Identical pattern for CNPG

Mount ImageVolume inside `extension_search_path`

You're done!*

*Modulo
`shared_preload_libraries`,
`ImageVolume` on pod start



FUN WITH DEPENDENCIES

DEPENDENCIES

Many extensions depend on shared libraries

Packagers manage them via normal means

Can Trunk format genericize?

Two approaches:

- 1. List dependencies, CLI installs from OS packaging**
- 2. Bundle dependencies**

Latter required for ImageVolume immutability

PACKAGED DEPENDENCIES

Package dependencies loaded by OS

ldconfig caches installed shared libraries

Perfect for packaged solutions

Infeasible for immutable systems

Cannot install in immutable directory

Can run ldconfig for new library in mutable directory

Cannot write to cache in immutable directory

BUNDLING DEPENDENCIES

Install dependent shared library with extension module

Compile Postgres with `RUNPATH` to that directory

No `ldconfig` caching required

`RUNPATH` supported by most OSes

lekensteyn.nl/rpath.html

RUNPATH

```
› chrpah $(which postgres)
/usr/lib/postgresql/bin/postgres:
RUNPATH=/usr/lib/postgresql/lib:/var/lib/postgresql/data/mod
```

```
› ls -l /var/lib/postgresql/data/mod
liburiparser.so.1
liburiparser.so.1.0.30
uri.so
```

RUNPATH

```
> chroot /var/lib/postgresql/data/mod/uri.so  
/var/lib/postgresql/data/mod/uri.so:  
RUNPATH=/usr/lib/postgresql/lib:/var/lib/postgresql/data/mod
```

```
> ldd /var/lib/postgresql/data/mod/uri.so  
linux-vdso.so.1  
liburiparser.so.1 => /var/lib/postgresql/data/mod/liburiparser.so.1  
libc.so.6 => /lib/aarch64-linux-gnu/libc.so.6  
/lib/ld-linux-aarch64.so.1
```

Loading with RUNPATH

```
postgres=# show dynamic_library_path;  
dynamic_library_path
```

```
-----  
$libdir /var/lib/postgresql/data/mod
```

```
postgres=# create extension uri;  
CREATE EXTENSION
```

```
postgres=# select 'https://example.com/'::uri;  
uri
```

```
-----  
https://example.com/
```

\$ORIGIN

Use rpath=\$ORIGIN

Points to same directory as library

Supported by all major Unixes

@loader_path on Darwin (macOS)

Not supported on Windows

\$ORIGIN

```
> chrpath --replace '$ORIGIN' /var/lib/postgresql/data/mod/uri.so
/var/lib/postgresql/data/mod/uri.so: new RUNPATH: $ORIGIN
```



```
> psql -c "select 'https://example.com/'::uri;"  
uri  
-----  
https://example.com/
```



```
> chrpath --replace 'NOTHING' /var/lib/postgresql/data/mod/uri.so
/var/lib/postgresql/data/mod/uri.so: new RUNPATH: NOTHING
```



```
> psql -c "select 'https://example.com/'::uri;"  
ERROR: could not load library "/var/lib/postgresql/data/mod/uri.so":  
liburiparser.so.1: cannot open shared object file:  
No such file or directory
```

BUNDLING DEPENDENCIES

Extract dependencies from OS packages

Put them in the Trunk lib directory

Compile extension with \$ORIGIN

```
make CFLAGS=' -Wl,-rpath,\$\$ORIGIN'
```

Postgres itself doesn't even need rpath

Success?

Descendant Dependencies

```
› chrpath /var/lib/postgresql/data/mod/http.so
http.so: RUNPATH=$ORIGIN:/usr/lib/postgresql/lib

› ldd /var/lib/postgresql/data/mod/http.so
linux-vdso.so.1
libcurl.so.4 => not found
libc.so.6 => /lib/aarch64-linux-gnu/libc.so.6

› scp dev:libcurl.so.4 /var/lib/postgresql/data/mod/
```

Descendant Dependencies

```
> ldd /var/lib/postgresql/data/mod/http.so
linux-vdso.so.1
libcurl.so.4 => /var/lib/postgresql/data/mod/libcurl.so.4
libc.so.6 => /lib/aarch64-linux-gnu/libc.so.6
/lib/ld-linux-aarch64.so.1
libnnghttp2.so.14 => not found
libidn2.so.0 => /lib/aarch64-linux-gnu/libidn2.so.0
librtmp.so.1 => not found
libssh.so.4 => not found
libpsl.so.5 => not found
libssl.so.3 => /lib/aarch64-linux-gnu/libssl.so.3
libcrypto.so.3 => /lib/aarch64-linux-gnu/libcrypto.so.3
libgssapi_krb5.so.2 => /lib/aarch64-linux-gnu/libgssapi_krb5.so.2
libldap.so.2 => not found
liblber.so.2 => not found
libzstd.so.1 => /lib/aarch64-linux-gnu/libzstd.so.1
libbrotlidec.so.1 => not found
```

Descendant Dependencies

```
> scp dev: libnghhttp2.so.14 /var/lib/postgresql/data/mod/  
> ldd /var/lib/postgresql/data/mod/http.so  
linux-vdso.so.1  
libcurl.so.4 => /var/lib/postgresql/data/mod/libcurl.so.4  
libc.so.6 => /lib/aarch64-linux-gnu/libc.so.6  
/lib/ld-linux-aarch64.so.1  
libnghhttp2.so.14 => not found
```

!!

```
libidn2.so.0 => /lib/aarch64-linux-gnu/libidn2.so.0  
librtmp.so.1 => not found  
libssh.so.4 => not found  
libpsl.so.5 => not found  
libssl.so.3 => /lib/aarch64-linux-gnu/libssl.so.3  
libcrypto.so.3 => /lib/aarch64-linux-gnu/libcrypto.so.3  
libgssapi_krb5.so.2 => /lib/aarch64-linux-gnu/libgssapi_krb5.so.2  
libldap.so.2 => not found  
liblber.so.2 => not found  
libzstd.so.1 => /lib/aarch64-linux-gnu/libzstd.so.1
```

WORKAROUND?

Works only for direct dependencies

Needs more experimentation

Maybe: Compile dependencies with own RUNPATHs

Cannot copy from OS packages

Unless they start using RUNPATH

Meanwhile, only bundle direct dependencies



FUTURE WORK

PROJECT STATUS

The vision remains in tact

Accessible, easy-install extensions everywhere

Close to completing v1 build SDK

Supports PGXS, pgrx

First deliverable: CLI

Work will inform RFCs

Update RFCs, discover patterns

STUFF I'D LIKE TO DO

Finish working out Trunk format and dependency patterns

Develop and submit `extension_search_path` patch

Submit ImageVolume feedback to Kubernetes

Allow runtime mounting!

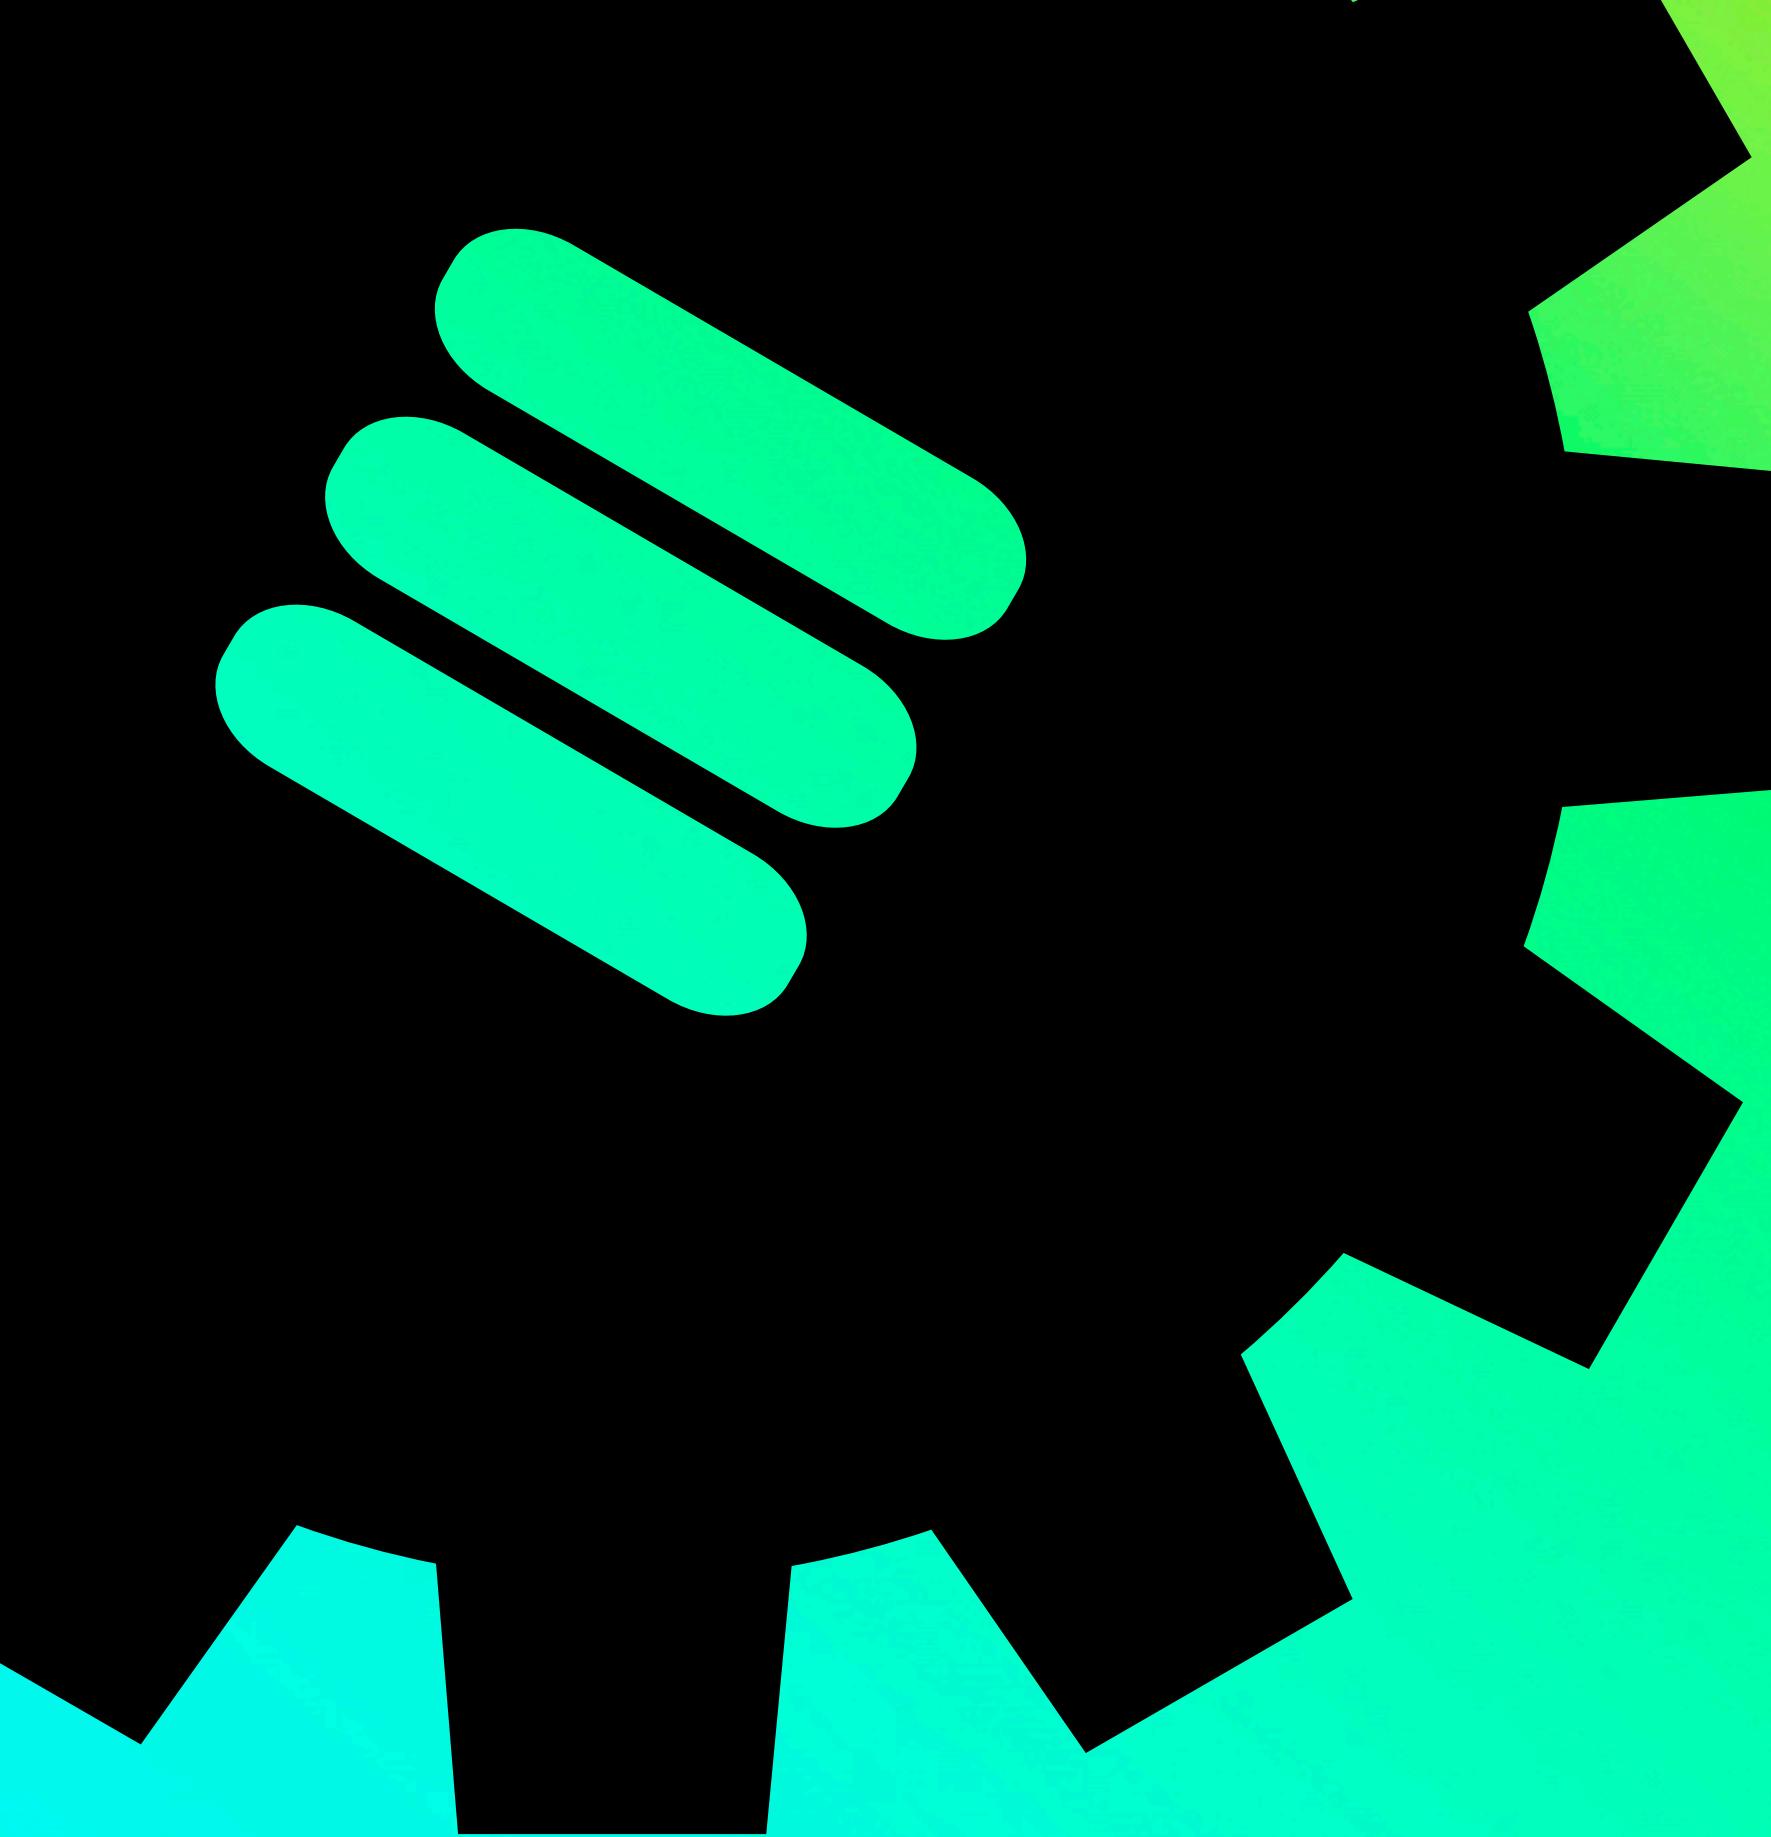
Start building and distributing OCI Trunk packages

Make pattern available for distributed registries

Make your own Trunk releases!

Hack fully-dynamic extension loading into CloudNativePG

LET'S TALK!



Thank You

Adventures in Extension Packaging
David Wheeler
PGXN, Tembo