

How can I deploy a
PostgreSQL extension in
Kubernetes without
breaking the immutability
of the container?





Improving the PostgreSQL Extensions Experience in Kubernetes with CloudNativePG

Postgres Extensions Ecosystem Virtual Mini-Summits

Wednesday 7 May, 2025

Gabriele Bartolini

VP, Chief Architect, Kubernetes at EDB



Gabriele Bartolini

VP, Chief Architect of Kubernetes at EDB

PostgreSQL user since ~2000

Ex 2ndQuadrant (co-founder)

PostgreSQL Contributor

DoK Ambassador

DevOps evangelist

Open source contributor

- Barman (2011)
- CloudNativePG (2022)



Blog: gabrielebartolini.it @_GBartolini_



#1 Contributors to



Creators of



CloudNativePG

enterprisedb.com

cloudnative-pg.io

postgresql.org



Agenda

- Setting the context
- Immutable Application Containers
- PostgreSQL Extensions in Kubernetes with CNPG
- The (immutable) future
- Conclusions





Setting the context

What is Kubernetes?

Greek for helmsman. Also known as **k8s**. URL: kubernetes.io

- Orchestration system for containers
- Automates deployment, administration and scaling of:
 - **infrastructure**
 - **cloud native applications** - also known as **workloads**
 - In Kubernetes, databases are workloads
- Open Source (Apache Licence 2.0) since 2014
 - v1.0 released in 2015
- Owned by the Cloud Native Computing Foundation (CNCF, cncf.io)
 - First project within the CNCF to become “Graduated”



CLOUD NATIVE
COMPUTING FOUNDATION



What is CloudNativePG?

- “Level 5”, Production ready
 - IBM Cloud Pak, Google Cloud, Azure, Akamai, Tembo, ...
- **CNCF Project (Sandbox)**
 - Apache License 2.0
 - Owned by the Community (CNCF)
 - Vendor neutral, openly governed, **always free**
 - Originally created by EDB (2019 by 2ndQuadrant)
- Multiple installation methods:
 - K8s manifests
 - Helm chart
 - OperatorHub.io (OLM)
- 4.0k commits
- 5.9k stars on Github (#1 Postgres operator)
- **~78M** downloads on Github (~3 years)



github.com/cloudnative-pg



FILTERS APPLIED (reset all):

PROJECT: CNCF



Application Definition & Image Build

dapr CNCF GRADUATED	HELM CNCF GRADUATED	ArtifactHUB CNCF INCUBATING	Backstage CNCF INCUBATING	Buildpacks.io CNCF INCUBATING	KubeVela CNCF INCUBATING	KubeVirt CNCF INCUBATING	OPERATOR FRAMEWORK CNCF INCUBATING	CARVEL MICROCKS	Devfile.io Nocloudhost	DevSpace Porter	ko radius	KONVEYOR porter	KUDO score	cloudevents sealer	NATS drasi	STRIMZI Praevent	Tremor
SHIPRIGHT	Knative	CloudEvents	Knative	Buildpacks.io	KubeVela	KubeVirt	OPERATOR FRAMEWORK	CARVEL	Devfile.io	DevSpace	ko	KONVEYOR	KUDO	cloudevents	NATS	STRIMZI	

argo CNCF GRADUATED	flux CNCF GRADUATED	keptn CNCF INCUBATING	OpenKruise CNCF INCUBATING
BRIGADE CNCF INCUBATING	OpenGDK CNCF INCUBATING	Pipeliner CNCF INCUBATING	werf CNCF INCUBATING
kube-burner CNCF INCUBATING			

TIKV CNCF GRADUATED	Vitess CNCF GRADUATED	TiKV CNCF GRADUATED	Vitess CNCF GRADUATED
Gemini CNCF INCUBATING	SchemaHero CNCF INCUBATING	TiKV CNCF GRADUATED	Vitess CNCF GRADUATED

Scheduling & Orchestration

KEDA CNCF GRADUATED	kubernetes CNCF GRADUATED	Crossplane CNCF INCUBATING	KARMADA CNCF INCUBATING	Knative CNCF INCUBATING	Kubeflow CNCF INCUBATING	VOLCANO CNCF INCUBATING	wasmcloud CNCF INCUBATING	Armarkat CNCF INCUBATING	capsule CNCF INCUBATING	ClusterNet CNCF INCUBATING	Clusterpedia CNCF INCUBATING	ERASER CNCF INCUBATING	envoy CNCF GRADUATED	CONTOUR CNCF INCUBATING	EMISSARY INGRESS CNCF INCUBATING	
KubeSlice CNCF INCUBATING	KubeStellar CNCF INCUBATING	Kured CNCF INCUBATING	Open Cluster Management CNCF INCUBATING	OPEN FUNCTION CNCF INCUBATING	Serverless Devs CNCF INCUBATING			FLUID CNCF INCUBATING	HAM CNCF INCUBATING	KCP CNCF INCUBATING	koordinator CNCF INCUBATING	kubers CNCF INCUBATING	BFE CNCF INCUBATING	LoxiLB CNCF INCUBATING	MetalLB CNCF INCUBATING	OpenELB CNCF INCUBATING

Istio CNCF GRADUATED	LINKERD CNCF GRADUATED	Aeraki	Heptio	KMESH	Kuma	Helm	Sermant	SIMP
-------------------------	---------------------------	--------	--------	-------	------	------	---------	------

CoreDNS CNCF GRADUATED	etcd CNCF GRADUATED	K8G8	Xline
---------------------------	------------------------	------	-------

gRPC CNCF INCUBATING	Connect
-------------------------	---------

Cloud Native Storage

FS	D	V	Carina
----	---	---	--------

KubeCon + CloudNativeCon Hong Kong June 10-11, 2024

KubeCon + CloudNativeCon Tokyo, Japan June 16-17, 2024

Native Network

Antria	CNI-Genie	FolksEdge	Rave
--------	-----------	-----------	------

What does cloud native means to us?

The three pillars of cloud native based on our experience since 2019

- **People and a “DevOps culture”**
 - Immutable Infrastructure
 - Infrastructure abstraction
 - Version Control
 - Declarative Configuration
 - Shift left on Security
 - Continuous Delivery
 - ...
- **Immutable Application Containers**
- **Kubernetes** as the standard way to orchestrate them through IaC





Immutable Application Containers

Mutable infrastructure

Prevalent mindset

- **Main artifacts: Packages**
- Long term infrastructure
- Changes happen as incremental updates on an existing infrastructure
 - dnf update or apt update
- Current state is represented by changes over time:
 - configuration management and operations
- Systems are healed after failures (*pets*)



Immutable infrastructure

Emerging mindset from DevOps and Platform Engineering

- **Main artifacts: OCI Container Images**
- Founded on Continuous Delivery (GitOps)
- Releasing a new version of an application means:
 - automatically build a new container image
 - publish it on a public/private registry
- Changes happen in atomic operations:
 - the new version image is pulled from the registry
 - the existing image is replaced “almost” instantaneously



Kubernetes: mutable or immutable?

Immutability requires a mindshift

- **There are no incremental updates**
- Data volumes are attached to containers (**stateful**)
- Containers are not changed (**immutable**)
 - A single change to the content, requires a new version of the image
- **Kubernetes works with immutable containers**
- **Security and Change Management Policies require this approach**
 - At any time in Kubernetes you know what software is running in your infrastructure, and it is in a version control system
 - Very important to control CVEs



The lifecycle of a container image

Simplified view



Immutable Application Containers

From day 1 of CloudNativePG we have placed this idea at the core

- The content of a running container must not change
 - File system is read-only to prevent any unauthorized change
 - “apt update” or “dnf update” are not possible
- Only a single application is running
 - PID 1 or entrypoint
- CloudNativePG and Immutable Application Containers:
 - Database files are in separate persistent volumes
 - PGDATA and optionally WAL files and tablespaces
 - Orchestrates “operand images” containing PostgreSQL binaries
 - The “instance manager” is the PID 1





	dependabot[bot] Bump sigstore/cosign-installer from 3.8.1 to 3.8.2 (#182)	1a8f19f · 4 hours ago	1,038 Commits
	.github Bump sigstore/cosign-installer from 3.8.1 to 3.8.2 (#182)	4 hours ago	
	Debian Automatic ClusterImageCatalog update	3 days ago	
	lib chore: update barman-cloud to version 3.13.3 (#177)	2 days ago	
	.gitignore chore: add PostgreSQL images 15beta (#18)	3 years ago	
	BUILD.md feat: add cosign to sign the images (#137)	4 months ago	
	CODEOWNERS Initial Commit	3 years ago	
	CODE_OF_CONDUCT.md Initial Commit	3 years ago	
	Dockerfile fix: Dockerfile to reduce vulnerabilities (#167)	3 weeks ago	
	GOVERNANCE.md Initial Commit	3 years ago	
	LICENSE Initial Commit	3 years ago	
	README.md feat: add cosign to sign the images (#137)	4 months ago	
	docker-bake.hcl chore(deps): update debian base images (#174)	last week	
	renovate.json Enable pinGitHubActionDigests for renovatebot (#165)	last week	

About

Operand images for CloudNativePG containing all community supported version PostgreSQL

[postgres](#) [postgresql](#) [container-images](#)
[cloudnativepg](#)

[Readme](#)

[Apache-2.0 license](#)

[Code of conduct](#)

[Activity](#)

[Custom properties](#)

[77 stars](#)

[6 watching](#)

[42 forks](#)

[Report repository](#)

Releases

No releases published

[Create a new release](#)

Packages 2

[postgresql](#)

[postgresql-testing](#)



Mmhh ...
What about PostgreSQL extensions?





PostgreSQL Extensions in Kubernetes with CloudNativePG

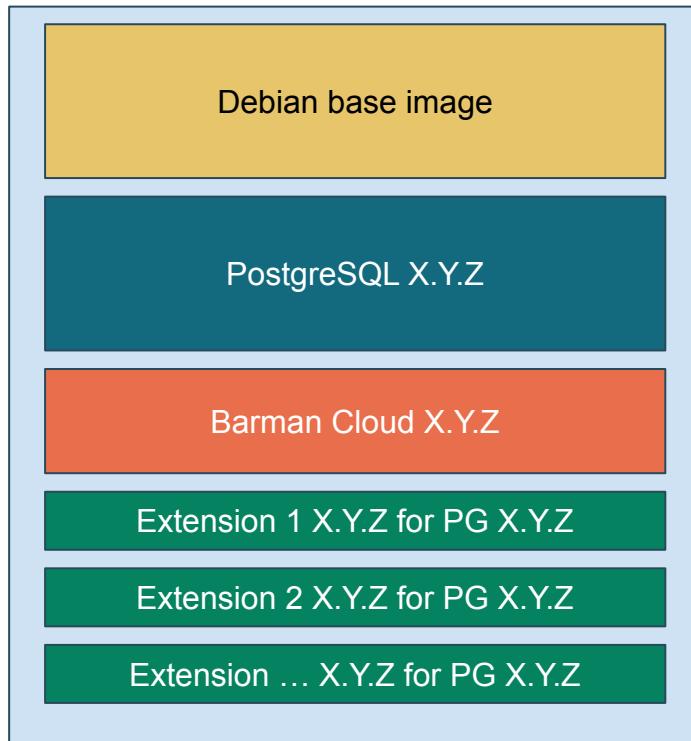
Possible scenarios for extensions

- Place them on-demand in a mutable volume
 - We immediately discarded that as it breaks immutability
- **Place them in the operand image**
 - Extensions must be included at build time in the operand image
 - A few approaches:
 - The “Spilo” way (images containing “the lot”)
 - **Create a few different flavours**
 - **Define a way to build custom images**
 - Container image requirements



PostgreSQL Operand Images now

How we distribute PostgreSQL container images in CloudNativePG



Every additional layer makes the image heavier

*Images are normally rebuilt every week in case
some element in the stack changes **

*Extensions are normally installed through
packages*



The problem with extensions

Remember: we cannot break immutability

- “Can I use extension WHATEVER in my CloudNativePG cluster?”
 - We add the extension to the container image
 - We could create a new flavor (e.g. PostGIS, TimeScale, ...)
 - Users ultimately need to build their own images
- Very large images!
 - System dependencies accentuate this problem
 - Increase the risk of CVEs being present in the system
 - Require more frequent updates (in terms of rebuild!)



Technology limitations

Both for PostgreSQL and Kubernetes

- PostgreSQL extensions must be in a single folder
 - e.g. /usr/share/postgresql/17/extension
 - It is not possible to define multiple locations
 - Until PostgreSQL 18 with “`extension_control_path`”!!!
- Kubernetes doesn't allow mount OCI artifacts as read-only volumes
 - Until Kubernetes 1.33 (released April 23, 2025)
 - Support for mounting images as volumes



Commit 4f7f7b0



petere and mattheusv committed 2 days ago · ✓ 7 / 10



extension_control_path

The new GUC extension_control_path specifies a path to look for extension control files. The default value is \$system, which looks in the compiled-in location, as before.

The path search uses the same code and works in the same way as dynamic_library_path.

Some use cases of this are: (1) testing extensions during package builds, (2) installing extensions outside security-restricted containers like Python.app (on macOS), (3) adding extensions to PostgreSQL running in a Kubernetes environment using operators such as CloudNativePG without having to rebuild the base image for each new extension.

There is also a tweak in Makefile.global so that it is possible to install extensions using PGXS into a different directory than the default, using 'make install prefix=/else/where'. This previously only worked when specifying the subdirectories, like 'make install datadir=/else/where/share pkglibdir=/else/where/lib', for purely implementation reasons. (Of course, without the path feature, installing elsewhere was rarely useful.)

Author: Peter Eisentraut <peter@eisentraut.org>

Co-authored-by: Matheus Alcantara <matheussssilv97@gmail.com>

Reviewed-by: David E. Wheeler <david@justattheory.com>

Reviewed-by: Gabriele Bartolini <gabriele.bartolini@enterprisedb.com>

Reviewed-by: Marco Nenciarini <marco.nenciarini@enterprisedb.com>

Reviewed-by: Niccolò Fei <niccolo.fei@enterprisedb.com>

Discussion: <https://www.postgresql.org/message-id/flat/E7C7BFFB-8857-48D4-A71F-88B359FADCFD@justattheory.com>





VolumeSource: OCI Artifact and/or Image #4639

[New issue](#)[Open](#)

sallyom opened on May 16, 2024 · edited by saschagrunert

Edits

...

Enhancement Description

- One-line enhancement description (can be used as a release note): VolumeSource: OCI Artifact or OCI Image
- Kubernetes Enhancement Proposal: <https://github.com/kubernetes/enhancements/tree/master/keps/sig-node/4639-oci-volume-source>
- Discussion Link:
- Primary contact (assignee): sgrunert@redhat.com, somalley@redhat.com
- Responsible SIGs: sig-node, sig-storage
- Enhancement target (which target equals to which milestone):
 - Alpha release target (x.y): 1.31
 - Beta release target (x.y): 1.33
 - Stable release target (x.y):

 Alpha KEP (k/enhancements) update PR(s):

- ↗ [KEP-4639: adding OCI VolumeSource #4642](#)

 Code (k/k) update PR(s):

- ↗ [\[KEP-4639\] Add OCI VolumeSource CRI API kubernetes#125659](#)
- ↗ [\[KEP-4639\] Update CRI API and workflow #4751](#)
- ↗ [\[KEP-4639\] Add ImageVolumeSource API kubernetes#125660](#)
- ↗ [\[KEP-4639\] Add ImageVolumeSource implementation kubernetes#125663](#)

 Nice to have:

- ↗ [\[KEP-4639\] Mention that fsGroupChangePolicy has no effect kubernetes#126281](#)
- ↗ [Fix runtime panic in imagevolume CanSupport method kubernetes#126323](#)
- ↗ [\[KEP-4639\] Add ImageVolumeSource node e2e tests kubernetes#126220](#)

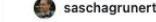
 Out of tree update PR(s):

- ↗ [Add OCI Volume Source support to crictl \[create|run\] --with-pull kubernetes-sigs/cri-tools#1464](#)

Assignees



sallyom



saschagrunert

Labels

[kind/design](#) [lead-opted-in](#) [sig/node](#)
[sig/storage](#) [stage/beta](#) [tracked/yes](#)

Type

No type

Projects

<input type="checkbox"/> 1.33 Enhancements Tracking
Status Tracked for Docs Freeze
<input type="checkbox"/> SIG Node 1.33 KEPs planning
Status Implemented

3 closed projects

Milestone



v1.33

No due date

Relationships

None yet





The (immutable) future

What's new in PostgreSQL 18

Introduce “extension_control_path”

- Path to search for extension control files (i.e. extensions)
- List of absolute directory paths
- Works with “dynamic_library_path” for additional libraries



PostgreSQL Operand Images in the future

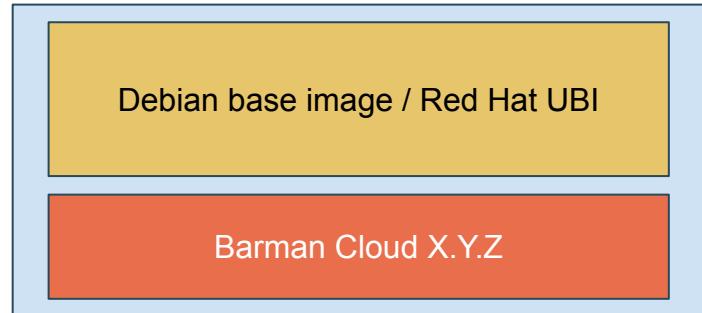
How we will distribute PostgreSQL and extensions in the future

PostgreSQL (minimal)



Every image will be independent and updated only when that software changes, or the underlying image (e.g. CVEs).

Barman Cloud Plugin sidecar



Any extension



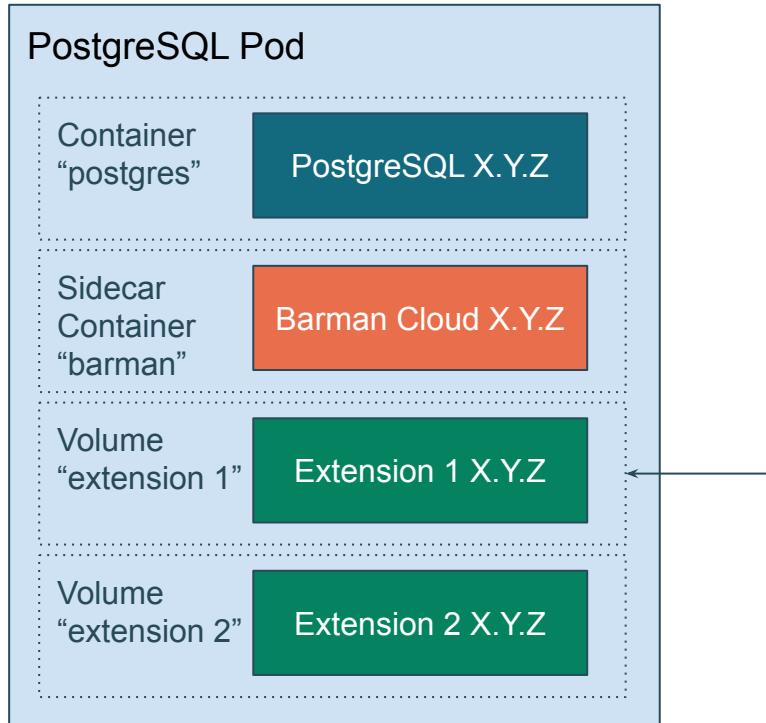
Note: no base image, just the actual files.

Obviously they must be coherent with the OS of the main PostgreSQL image (e.g. if Debian, extension should be built from the Debian package content or compiled)



What will change at runtime (without breaking immutability)

Extensions can be added at any time, without having to prebuild images with all needed combinations



*extension_control_path could list
all the extension volumes loaded via
Kubernetes VolumeSource*

*This could be
/extensions/pgvector/extension*



Fail-fast pilot project

Introduce extensions management in CloudNativePG

- Exploratory patch in the CloudNativePG operator:
 - Declarative management of extensions for a “Cluster” resource
 - Introduce the “.spec.postgresql.extensions” stanza
- PostgreSQL image with “extension_control_path”
 - Built from the “postgres-trunk-containers” project
- Bare container images for a few extensions, including:
 - pgvector
 - postgis
- *Many thanks to Niccolò Fei and Marco Nenciarini for working on the pilot*



NiccoloFei requested your review on this pull request.

[Add your review](#)

feat: add support for configuring PostgreSQL extensions via Image Volume #6546

[Edit](#) [Code](#)[Draft](#)NiccoloFei wants to merge 7 commits into `cloudnative-pg:main` from `EnterpriseDB:dev/5645`[Conversation 2](#)[Commits 7](#)[Checks 27](#)[Files changed 10](#)[+270 -0](#)

NiccoloFei commented on Jan 9 · edited

Member

...

Add support to allow setting up PostgreSQL extensions via container images, taking advance of the `ImageVolume` `VolumeSource` of Kubernetes and the new proposed GUC of PostgreSQL `extension_control_path`.

Introduce a new API field `Extensions` in the `cluster.Spec.PostgresConfiguration` stanza, which takes an array of PostgreSQL extensions which should be added to the Cluster.

For more information, see:

- <https://github.com/kubernetes/enhancements/tree/master/keps/sig-node/4639-oci-volume-source>
- <https://kubernetes.io/docs/tasks/configure-pod-container/image-volumes/>
- <https://commitfest.postgresql.org/50/4913/>
- <https://justatheory.com/2024/11/rfc-extension-packaging-lookup/>



NiccoloFei requested review from `jsilvela` and `cloudnative-pg/maintainers` as code owners 4 months ago

Reviewers

jsilvela



maintainers



At least 2 approving reviews are required to merge this pull request.

Assignees

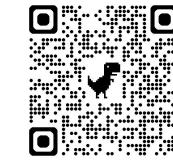
No one—assign yourself

**Labels**

do not backport

**Projects**

None yet

**Milestone**

Our goal for CloudNativePG 1.27

```
apiVersion: postgresql.cnpq.io/v1
kind: Cluster
metadata:
  name: postgresql-with-extensions
spec:
  instances: 1
  imageName: ghcr.io/cloudnative-pg/postgresql-trunk:18-devel
  postgresql:
    extensions:
      - name: pgvector
        image:
          reference: ghcr.io/cloudnative-pg/pgvector-18-testing:latest
  storage:
    storageClass: standard
    size: 1Gi
```



What will happen under the hood

- CNPG triggers a rolling update, starting with replicas
 - *Short downtime required*
- Each referenced image is mounted as a read-only volume via *ImageVolumes*
 - The *pgvector-18-testing:latest* image is mounted on */extensions/pgvector*.
- CNPG updates:
 - *dynamic_library_path* to include the */extensions/pgvector/lib*
 - *extension_control_path* to include */extensions/pgvector/share*

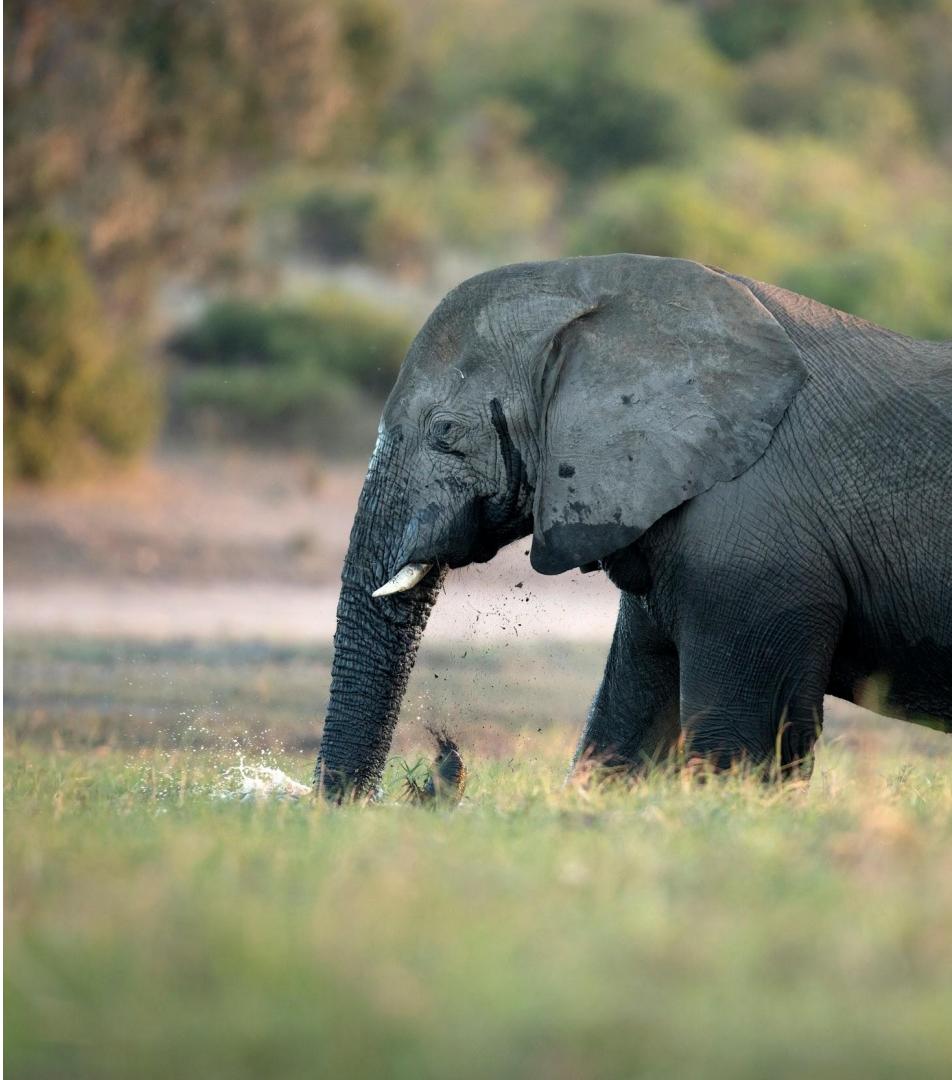


Container Images as first class artifacts

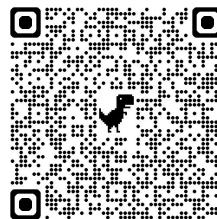
Container Images should become a standard distribution artifact for extensions, like packages

- Currently, extensions are primarily distributed as packages
- We should distribute them as OCI images
 - Epical change for PostgreSQL





Conclusions



Suggested reading

The Immutable Future of PostgreSQL Extensions in Kubernetes with CloudNativePG

3 March 2025 · 8 mins

kubernetes k8s cloudnativepg cnpg postgresql postgres dok data on kubernetes extensions
Container images sbom pgvector imagevolume extension_control_path



Key take aways

The future for PostgreSQL extensions in Kubernetes is immutable

- Container images should be seen as a standard artifact nowadays
 - Challenge the status-quo of Postgres extensions distribution
- An ecosystem based on immutable container images will bring:
 - Reduced operational complexity, easier testing and validation of extensions
 - Lighter and safer images, with less frequent need for updates and bandwidth usage
 - Full Independence for extension projects
 - Each extension will ideally build their own images and publish them
 - No need to centralize build of extensions and prebuild combinations
- Refrain from saying “This is a limitation of Kubernetes for Postgres”
 - We need to consider all the organisational benefits brought by the stack



Questions?

gabrielebartolini.it @_GBartolini_



Creators of  **CloudNativePG**

enterprisedb.com

cloudnative-pg.io

cloudnativepg@mastodon.social

cloudnativepg.bsky.social

linkedin.com/company/cloudnative-pg/